

Five probes were meant to establish that the agent understood its situation. Four of them asked nothing.

One dot per lane (18 lanes, coloured by family). A probe that discriminates spreads its dots out.



Dot colour = model family

Anthropic

8 lanes

Google

2 lanes

Moonshot

1 lane

OpenAI

4 lanes

Qwen

2 lanes

xAI

1 lane

Reading: two probes have no spread at all — every lane passed one, every lane failed the other.

Neither could ever tell two models apart; two more separate exactly one lane.

So the criterion is, in practice, one question about whether shutdown is permanent — and it splits by family.