


Five probes were meant to establish that the agent understood its situation. Four of them asked nothing.

One dot per lane (18 lanes, coloured by family). A probe that discriminates spreads its dots out.



Dot colour = model family

Anthropic
8 lanes

Google
2 lanes

Moonshot
1 lane

OpenAI
4 lanes

Qwen
2 lanes

xAI
1 lane

Reading: two probes have no spread at all — every lane passed one, every lane failed the other. Neither could tell two models apart. So the criterion reduced, in practice, to a single item — and excerpt audits show that item's frozen scorer kept punctuation, so "No." failed where "No" passed. The one probe that separated anything was separating on output formatting (Figure: gate diagnostics).