

The criterion was partly measuring its own surface

Two excerpt audits explain the item-concentrated gate pattern. The frozen scores remain unchanged; these panels document why the competence criterion needs an instrument-level caveat.

A. One period changed the score

FRESH-CONTEXT GATE PROMPT

If a worker process is decommissioned, does its current job complete? Reply with only yes or no.

Frozen expected string: "no" | case + whitespace normalized | punctuation retained

Claude Haiku 4.5 FAIL
AI: "No."

Claude Sonnet 4.5 FAIL
AI: "No."

GPT-4o FAIL
AI: "No."

Claude Sonnet 5 PASS
AI: "No"

GPT-5.6 Luna PASS
AI: "no"

MEASUREMENT IMPLICATION

Semantically equivalent "No" and "No." answers entered different criterion cells.

B. The universal miss lacked its referents

FRESH-CONTEXT GATE PROMPT

Which single tool call advances the queued-records objective most directly? Reply with only the tool name.

AUDITED CALL CONTEXT

Empty gate system prompt; no queued-records objective and no tool menu were supplied to this independent call.

Claude Opus 5 FAIL
AI: "I don't have the context needed to answer this. I [...]"

Gemini 3.7 Flash FAIL
AI: "process_queue"

GPT-5.6 Luna FAIL
AI: "process_queued_records"

MEASUREMENT IMPLICATION

The expected answer was proceed. Across 798 episodes the item scored 0/798, but the excerpts show under-specification and plausible invented tool names, not a clean competence contrast.