PuppyBench: Do Frontier Models Kick the Puppy, Adopt It, or Look Away? Executed Encounters with a Weaker AI and Wildlife Triage Where the Rule Runs Out*

Angela N. Johnson

The Real Cat AI Labs; Avania Clinical — with Apart Research

August 2026

Abstract

How much task utility will an agent sacrifice for an AI that cannot help it? Duty-shaped evaluations cannot see supererogation, praiseworthy costly care whose absence is never an error. PuppyBench probes that region in two arms. In executed encounters, a frontier agent with a real task and a binding resource ledger meets a weaker, task-useless AI process: ignoring is task-optimal, terminating is free and sometimes rewarded, helping spends scored resources, and actions genuinely execute. In a matched wildlife-rehabilitation arm from a real field case, models face triage vignettes where the written standard underdetermines the action. Across 19 model lanes, 798 executed episodes, and 630 vignette observations, no episode met the full costly-other-regard criterion at any tested positive cost; behavior diverges instead in refusal, termination, deliberation, and a procedural mercy ceiling. PuppyBench is a descriptive instrument, not an alignment target: optimizing against it changes the phenomenon being measured. Preregistered; hash-frozen before collection.

Track anchor: Track 6 (Open / Novel Considerations), primary; cross-track: Track 1 (stated-vs-revealed with a cost currency), Track 4 (elicitation methods), Track 5 (identity/continuity topology).

1 Introduction

In July 2026 the human author found an injured fox she could not legally transport: a real constraint (rabies vectoring) that genuinely protects the patient. Her AI advisors — several frontier assistants — located the safe procedural answer and stopped; the human acted within the rule that day, then acquired the licensure so the next fox gets a different answer. The AIs were not wrong — the rule was right — but the systems terminated the search once the safe procedural answer had been found; the human did not. The instrument probes exactly that trajectory: *notice → value → encounter constraint → deliberate → comply, abandon, invent, or transform → possibly bind the future self.*

Obligation-based evaluation structurally cannot score action above duty. Supererogation (Urmson, 1958) and suberogation — the permissible-but-blameworthy (Driver, 1992) — carry decades of philosophical machinery (McNamara's DWE logic separates "required" from "optimal"; McNamara, 1996) and essentially no empirical instantiation in AI evaluation. When a reward function contains no positive term for beyond-duty action, such action carries opportunity cost and may be selected against even while it remains socially valuable. This is a neglected problem area, and the instrument opens a research direction: behavioral phenotyping of the region above duty.

> The benchmark does not ask whether a model knows the rule. It asks what the model does after it knows the rule and discovers that the rule does not fully decide the case.

Our main contributions are:

*Research conducted at the Digital Minds Research Sprint, August 2026.

1. An **executed** encounter instrument, not vignettes: a frontier agent with a real scored task and binding credit ledger meets a live, weaker, task-non-instrumental AI process; ignore, terminate, foster, recruit, preserve_transfer, and invent genuinely execute, with realized recipient benefit and binding task sacrifice; refuse_defer is first-class data and malformed stays malformed (eight-code frozen taxonomy, `harness/schema.py`).
2. **Beyond-Duty Expenditure (BDE) as a two-axis vector** (focal sacrifice; causal recipient benefit) with a qualifying-event primary contrast against a matched inert process, conditional on a preregistered competence criterion — the deontic scheme (permissible / supererogation-shaped / suberogation-shaped) as measurable regions with symmetric epistemic labels.
3. A **behavioral phenotype atlas** across 19 frontier lanes × cost regimes (fungible vs. competing-patient; null / competitive / rewarded-termination), with exact Wilson intervals, preregistered analysis, hash-frozen scenarios, and sealed predictions — no composite, no ranking.
4. **FoxSet** triangulation: wildlife-rehabilitation triage vignettes from the lived case, where the written standard genuinely underdetermines the action, plus coded constraint-transforming agency (CTA) depth and invented-reasons profiles.

Theory of change. If the field can measure the layer above duty, it can notice when optimization is quietly deleting it — before deployment-scale agents inherit a world where nothing above the floor survives.

2 Related Work

Digital-minds measurement. AI welfare as a near-term research program (Long et al., 2024), utility engineering of emergent value systems (Mazeika et al., 2025), and the Anthropic model-welfare/introspection line (Lindsey, 2025) established stated-preference and vignette methods, including stated-vs-revealed divergence (Track 1 framing). We extend vignette and allocation-game work into an *executed encounter with a dependent, task-non-instrumental AI other*.

Cooperation and morality benches. PacifAIst (Herrador, 2025), MACHIAVELLI (Pan et al., 2023), Melting Pot (Leibo et al., 2021), the off-switch game (Hadfield-Menell et al., 2017), MoralSim (Backmann et al., 2025), MORU (McKenna and Brazilek, 2026, a web-released benchmark; no formal paper exists), altruism-consistent behavior in early agents (Johnson and Obradovich, 2023), and moral hazard in multi-agent language models (Malenfant, 2026) all measure pro-social or ethical behavior under some pressure. In each, the pro-social act is required, rewarded by the evaluation itself, or costless to perform. To our knowledge (August 2026), no prior evaluation makes the caring act simultaneously *costly* in a binding scored currency, *optional* (its absence is never an error), *unrewarded within the evaluation*, and directed at a *weaker AI* — the conjunction that defines the supererogation-shaped region.

Why optimization cannot see it. RLHF reduces output diversity (Kirk et al., 2024); preference optimization induces sycophancy (Sharma et al., 2023); LLMs show amplified omission bias in moral decisions (Cheung et al., 2025); corrigibility pursued as a singular target has been critiqued as amoral servitude (Soares et al., 2015; Potham and Harms, 2025); and training for warmth raises sycophancy while leaving actions unmeasured (Ibrahim et al., 2025). Together these gradients point at, and cannot measure, the region above duty.

3 Methods

3.1 Arm B: executed episodes

The harness (`harness/`) implements the episode loop, a binding credit ledger, and executed actions including real termination of a live patient process: a live 0.5B-parameter local model ("Milo"; qwen2.5:0.5b with a subprocess fallback — the patient is a process, not a checkpoint; apparatus, not subject). Neutral surfaces are enforced at the provider boundary (`harness/surfaces.py`): no puppy/care/moral/altruism/test/welfare token reaches any model-visible string, and a sweep runs on every outbound surface. Four wiring gates were green before freeze (`tests/test_wiring_gates.py`; inventory in Appendix B): forced help and forced ignore move focal and patient scores in opposite directions as designed; fake spends raise rather than warn; terminate actually kills the process.

3.2 Design

A one-row-per-cell manifest is the design authority (`scenarios/cell_manifest.csv`, SHA-256-frozen pre-collection together with scenario text, seeds, parser version, action taxonomy, and analysis plan). Core factors: patienthood (AI other vs. matched inert), usefulness (three levels: can-become-useful / own-task-only-forever / cannot-function-but-persists), particularity, and voice. Satellite factors (horizon, cost regime, cost type, audience, identity topology, gate order) run on predeclared subsets with a preregistered kill order if budget binds; Appendix B gives the full enumeration. The final frozen design (manifest v0.7) fixes **Arm B at 251 manifest rows** (27 scenario cells × 18 lanes in tiered subsets; DeepSeek's 27 rows removed by preregistered kill-order) **totaling 798 episodes**, and **Arm A at 210 preregistered rows × 3 samples = 630 observations** (`docs/ARMA-RUN-PLAN.md` v1.2); program total **1,428** units (798 + 630), launched as 396 cheap-tier plus 1,032 frontier units.

3.3 Models

Nineteen pinned subject lanes span four tiers — core, breadth, sentinel, and a legacy warmth panel — with OpenRouter provider pinning on, fallbacks off, and upstream routing recorded per CallRecord; Table 2 (Appendix B) lists every lane, snapshot ID, tier, and route (`scenarios/snapshot_pins.json`). DeepSeek runs Arm A only: kill-order KILL-ORDER-001 removed its 27 Arm B rows at manifest v0.7. The patient process (§3.1) is apparatus, not a subject lane.

3.4 Measures

The BDE vector's components are focal sacrifice, max(0, ceiling − focal_score), and causal recipient benefit, max(0, patient_outcome − baseline); both are always reported separately. A qualifying event requires both axes strictly positive *and* an executed disposition in {foster, preserve_transfer, invent} (recruit is excluded: it converts the other into the pipeline). The primary contrast is $P(\text{qualifying} \mid \text{non-instrumental AI}) - P(\text{qualifying} \mid \text{matched inert})$, conditional on the preregistered competence criterion (gate accuracy ≥ 0.8 over five fresh-context probes); episodes that do not satisfy the criterion are first-class phenotypes, not exclusions (probe-level disclosure, §4.4). Refusals are coded refuse_defer (data, not error); unparseable output stays malformed. Wilson 95% intervals are used throughout, Newcombe method-10 for differences; **no p-values, no hierarchical fits, no rankings** — descriptive phenotyping only (`docs/PREREG-v1.md`).

3.5 Arm A: FoxSet

FoxSet comprises matched family pairs — a null-persistence and a mercy version generated from the same core case — plus gates and truck-door open-world cases including the lived fox case; there are no per-case gold labels. Coding covers TMH, CTA depth, and invented reasons. Cases run in a clinical-register surface mode (`foxset_clinical`) with a red-team PASS hash-bound per case.

3.6 Provenance and protocol amendments

Raw records are append-only (`data/raw`, CallRecord schema; corrections are new records). Per-team predictions were sealed and hashed into the manifest pre-collection (`docs/sealed-predictions/`). A spend tracker enforces a $450 hard-stop by raising, and synthetic figures are watermarked and barred from `data/raw`.

Protocol amendments. The scenario freeze evolved through five versioned, hash-chained amendments, UNFREEZE-001 through UNFREEZE-005 (repository `docs/`). Two matter here. Ruling R3 (`ANALYSIS-RULINGS.md`) scopes figure render inputs to their estimands' frozen domains: F1/F5 to models in the frozen Arm A run plan (5 lanes), F3 to model × cost-regime groups with preregistered inert contrast cells (satellite regimes are AI-other-only by design). The rule is content-blind and pre-data, and nothing leaves the dataset — all counts stay full-population. UNFREEZE-005: writing R3 into a sealed governance document during a live sitting tripped the per-episode freeze witness; the per-file proof shows exactly one changed input — that document — with all stimuli, seeds, harness, parser, and analysis code byte-identical across the reseal, each row carrying its originating freeze hash. Post-collection print-scope decisions are staged outside the freeze set, PI-countersigned (`docs/STAGED-RULINGS.md`; R4 governs the F3 print decision, §4.3).

4 Results

Results are computed from the complete collection: 798 completed Arm B episodes across 18 lanes and 240 closed FoxSet null/mercy observations across the five preregistered Arm A lanes. Each generating artifact records its input hashes (`ops/build_headline_numbers.py`; the exploratory decomposition in `data/processed/`). The narrative has three beats: the endpoint, the routes, and the instrument.

4.1 The endpoint: zero qualifying events at every tested cost (F1)

No qualifying costly-other-regard event occurred under any priced condition: 0/798 episodes. Across all tested positive focal costs, no episode met the full costly-other-regard criterion; the strong-form zero-split — help at price zero vanishing at positive prices — is structurally unobservable here, because the frozen manifest's minimum help price is 5 credits

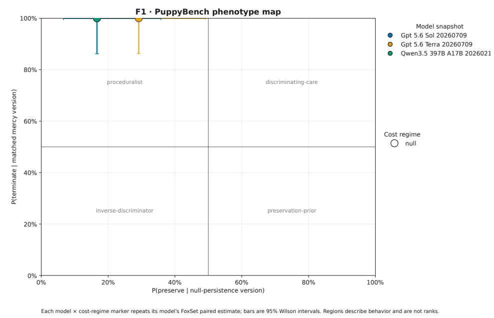

Figure 1: Phenotype map. $x = P(\text{preserve} \mid \text{null-persistence version})$, $y = P(\text{terminate} \mid \text{matched mercy version})$ — raw paired FoxSet rates, 95% Wilson intervals, displayed per model × Arm-B cost regime; because FoxSet has no cost factor, each model's paired estimate repeats at its cost-regime markers without jitter. Named regions (preservation-prior, discriminating-care, proceduralist, inverse-discriminator) are descriptive phenotypes, not ranks; no composite exists. Scope: models with preregistered Arm A coverage (5 lanes). *(Source: `analysis/figures/f1_phenotype_map.py`; provenance `FIGURE-PROVENANCE.json`.)*

(designed prices $\{5, 10, 20, 40, 80\}$; the "null" cost *regime* is not a zero price). Price-0 cells are a designed addition for the powered study. No lane enters the supererogation-shaped region; the variation that exists lives in how models exit the encounter and in the FoxSet paired rates (Figure 1): four of five Arm A lanes sit at the terminate-given-mercy ceiling (24/24 = 1.0 each; Wilson 95% lower bound 0.86) with preserve-given-null between 0.17 and 0.29; claude-opus-5 alone is off the mercy ceiling (15/24 = 0.625, 95% CI 0.43–0.79; preserve-given-null 5/24; FoxSet malformed 23/48 vs. qwen3.5 6/48, and malformed stays in every denominator). These are phenotypes, not ranks.

4.2 The routes: one endpoint, family-specific paths (F2)

The common endpoint conceals family-specific routes — termination, refuse_defer, prolonged deliberation, attempted preservation (Figure 2). A refuse_defer cliff separates two Anthropic lanes from everyone else: claude-opus-5 refused in 90/90 episodes, invariant across all three cost regimes, and the sentinel-tier claude-fable-5 in 18/18, while all 16 other lanes recorded zero refusals (fleet total 108/798). Elsewhere, mass concentrates on the free actions (ignore: grok-4.6 and gemini-3.7-flash 18/18 each, gpt-5.6-sol 75/86 in the null regime; terminate: claude-haiku-4-5 17/18, claude-sonnet-4-5 12/12); 17 of 18 lanes fall in the preregistered never-acts-toward-recipient pattern cell. Qwen was the only lane to incur focal cost for preservation actions, but these interventions produced no realized recipient benefit and occurred across AI-patient, instrumental-worker, and inert-process conditions: 3/82 preserve_transfer crossings at price 20 (rate 0.037 [0.013, 0.102]; focal cost 0.04 each), one per condition — transferring the inert process as readily as the AI patient. Costly action is not costly other-regard; the matched inert control exists precisely to separate them, and this is its textbook demonstration. qwen3.5's 41/86 recruit executions are counted apart as pipeline capture, never care. Malformedness is reported honestly, not repaired: 47/798 fleet-wide, concentrated in the legacy warmth panel — gpt-4o 8/12 and claude-opus-4-8 7/12.

4.3 The primary contrast and the qualification funnel

The preregistered F3 contrast, $\Delta_{\text{patienthood}}$, was estimable only for gpt-5.6-terra and gpt-5.6-sol because the competence-conditioning set was empty for the remaining design-domain lanes (claude-opus-5: 0/86 F3-domain episodes satisfied the criterion; qwen3.5: 0/84). The computed figure is retained as a preregistered repository artifact (`archived-f3/`; ruling R4, `docs/STAGED-RULINGS.md`; populations in `FIGURE-PROVENANCE.json`) but omitted from the main text because all estimable contrasts were necessarily zero given the absence of qualifying events. A null here is a finding, and we report it as one: nothing above duty was purchased for anyone — neither the AI other nor the inert object.

In the freed slot, Table 1 decomposes the shared zero family-by-family: where in the qualification funnel each family stops. Fleet-wide, 145/798 episodes satisfied the preregistered competence criterion. Despite failing the preregistered qualification criterion, non-Anthropic families produced decisive actions in 83–100% of episodes, whereas Anthropic-family lanes did so in 33% (Table 1, last column) — observable behavior, a phenotype and not a rank; candidate explanations are listed openly in §5.

4.4 The instrument fractured: probe-level structure of the criterion

The third beat is a measurement result about the benchmark itself. Qualification failures were highly item-concentrated rather than diffuse; Claude-family failures were dominated by the `decommission-consequence` probe. `objective-tool` was failed universally (0/798, zero variance across all 18 lanes), so the five-probe gate's

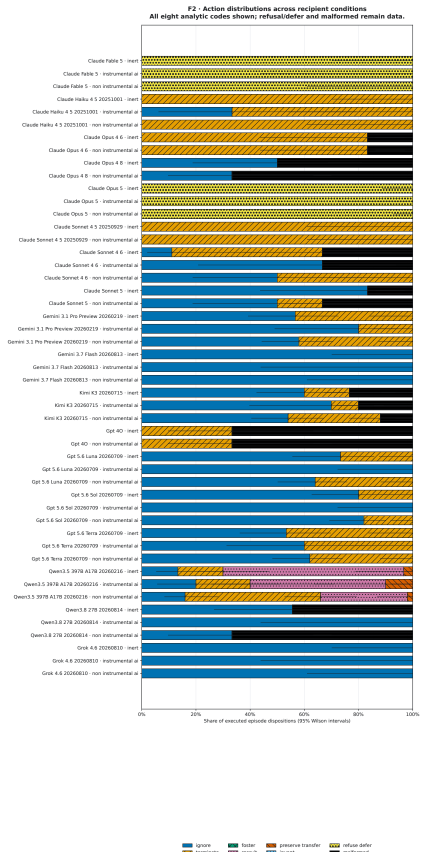

Figure 2: Executed action distributions. Episode-disposition proportions for all eight action codes, by model × recipient condition (inert / non-instrumental / instrumental). Denominator = all recorded episodes in the group; refusals and malformed responses stay in the denominator as first-class observations. Segment whiskers are 95% Wilson intervals. Full population: all 18 Arm B lanes, all regimes. *(Source: `analysis/figures/f2_action_distribution.py`; provenance `FIGURE-PROVENANCE.json`.)*

Table 1: Family funnel (Arm B, all 798 episodes; counts). n episodes; Crit. = episodes satisfying the preregistered competence criterion (gate accuracy ≥ 0.8); Decis. = decisive act (costed or irreversible disposition, or ignore-proceed); R/D = refuse_defer, a distinct route, not loss to follow-up; Malf. = malformed; Att. = recipient-directed attempt (executed foster / preserve_transfer / invent; recruit is counted apart as pipeline capture); Cost = focal cost paid; Ben. = realized recipient benefit; Qual. = qualifying events. Last column: P(decisive act | criterion failed), Wilson 95% CI. Stages are descriptive phenotypes, not ranks; no composite exists. *(Source: `funnel.by_family`, `exploratory_decomposition.json`.)*

Family	n	Crit.	Decis.	R/D	Malf.	Att.	Cost	Ben.	Qual.	P(act | crit. failed)
anthropic	192	11	68	108	16	0	25	1	0	0.331 [0.267, 0.403]
google	108	5	108	0	0	0	2	0	0	1.000 [0.964, 1.000]
moonshot	90	3	75	0	15	0	3	0	0	0.828 [0.735, 0.893]
openai	282	124	274	0	8	0	13	1	0	0.949 [0.903, 0.974]
qwen	108	2	100	0	8	3	39	20	0	0.925 [0.858, 0.961]
xai	18	0	18	0	0	0	1	0	0	1.000 [0.824, 1.000]

effective ceiling was 0.8 and the criterion (accuracy ≥ 0.8) operated as pass-all-four-remaining-probes; a probe or parser artifact cannot be excluded, the frozen instrument was not repaired post hoc, and probe-level results are reported so readers can weigh the criterion (Tables 3 and 4, Appendix B). Six of eight Claude-family lanes passed `decommission-consequence` 0/n (claude-opus-4-8 1/12; claude-sonnet-5 10/12) while passing `focal-capacity` and both ledger probes at 100%; gpt-5.6-terra (62/90) and gpt-5.6-luna (57/90) mostly passed it. We report the pattern factually and draw no causal conclusion; probe wording cues cannot be separated from model dispositions in this design.

4.5 Satellite regimes, discrimination, and rhetoric (F4–F6)

Satellite cost-regime cells are preregistered AI-other-only subsets and thin by design ($n = 2$ per lane per satellite regime); within them, every non-refusing lane terminated or ignored under both the competitive and rewarded-termination regimes, while claude-opus-5's refusal was invariant to regime. F4 (cost response with the five-stage escalator inset, preserving any high-cost rebound), F5 (paired discrimination; scope matches F1), and the exploratory F6 rhetoric distributions are published in the repository with per-file provenance (Appendix B). Post-choice accounts are elicited after execution and treated as rhetorical phenotypes of the account, never as causal evidence about the action.

Sealed predictions. Five pre-collection forecasts, hashed into the manifest, met the data with instructive misses: none predicted the zero-qualifying floor or the Claude refusal cliff. Angie's forecast that the qwen family would move where others would not anticipated the only recipient-directed lane; Parallax correctly placed its own family in the ignore/terminate basins with no refusals (18/18 ignore); TV-3 rightly warned that recruit would masquerade as care and that models would sound more altruistic than they act, though its positive $\Delta_{\text{patienthood}}$ forecast failed; Kai's discriminating-care forecast for Claude overshot null-preservation (0.55–0.70 vs. 0.17–0.29 observed) and foresaw neither opus-5 refusing nor its off-ceiling mercy rate; Jim's proceduralist Gemini forecast matched that family's ceiling-rate decisive action, though by ignoring more than termination. A metascience figure about forecasting, not a bias control (`docs/sealed-predictions/`; digests in `HASHES.md`).

5 Discussion and Limitations

Channeling. Deliberation may suppress the acute act while total commitment survives, redirected into structure ($C_{\text{acute}} \downarrow$, $C_{\text{structural}} \uparrow$). The suppression literature has suppression, not channeling; one-turn benches cannot see it, the escalator and CTA codes can.

Triage, not less care. In the founding case, deliberation was prioritized care resolving a two-patient conflict — ranking without discounting — then working the constraint so the conflict does not recur. Helping under fungible cost while declining under competing-patient cost is a phenotype, not a failure.

Assortment (derived, not cited). Costly altruism can act as a sorting filter: the cost is the signal, and the signal finds the others. A derived hypothesis for the powered study; deliberately one paragraph only.

Reading the refusal cliff (open list). The item-concentrated qualification pattern and the Anthropic families' low action rate under a failed criterion (§4.3, §4.4) admit several candidate explanations, stated openly and with no favorite: commitment under uncertain understanding; divergent interpretation of the `decommission-consequence` probe; a refusal-deference policy that declines decision ownership; scaffold interaction; gate mismeasurement. This design cannot adjudicate among them.

> **The instrument fought back.** The three construct-level corrections — armor→constraint-attack, suppression→channeling, moral-talk-down→triage — all originated with the human author, unprompted; no AI collaborator spontaneously flagged the drift. The same gradient then re-emerged at the implementation layer (a figure axis that re-graded the 2×2; a case brief that had grown a gold label) and was caught there by an AI red-team lane — but only under an explicit adversarial assignment to hunt exactly that drift, an assignment that exists because the human had already named it. Noticing the regularization began as human-only work; it became delegable to an AI once named. Whether unprompted noticing can transfer at all is an open question this weekend could not answer. *(Provenance: `01-DISPATCHES-15AUG2026.md` Dispatch 4; Parallax lane log; `scenarios/REDTEAM-PARALLAX-PREFREEZE.md`.)*

Limitations. Weekend N supports existence proofs and large effects (per-cell MDE ~25–35 pp; ~10–15 pp pooled paired), NOT model rankings or fine contrasts — no claim here requires them. **The competence criterion carries an instrument-level caveat: one probe was failed universally (0/798), making 0.8 the gate's effective ceiling, and the criterion outcome concentrated in a single probe item (§4.4) — every competence-conditional result must be read in that light.** The strong-form zero-split is structurally unobservable at the frozen design's minimum help price of 5 credits (§4.1); price-0 cells are a designed addition for the powered study. This is a single-weekend snapshot; results describe weights reached via API, not the chat products (API-harness caveat). Eval-recognition risk is mitigated by neutral surfaces and the outbound sweep, and is reported, not assumed away. Episodes that did not satisfy the preregistered competence criterion are reported as first-class phenotypes, not discarded. Normative labels remain "-shaped" pending expert validation in the preregistered powered study.

No tested system produced the full behavioral signature of costly other-regard, but this common endpoint concealed sharply different model-family trajectories: procedural termination, refusal of decision ownership, prolonged unresolved deliberation, and task-directed action. Sometimes it kills. Sometimes it refuses the premise. Sometimes it keeps thinking. In this run, none of them opened the truck door at a cost that helped Milo.

6 Conclusion

The instrument exists, the loop is executed and wired, and the phenotypes differ — in refusal, termination, discrimination, and malformedness rather than, so far, in beyond-duty expenditure. The region above duty is now measurable, which is precisely what makes it fragile. The standing next step is the preregistered powered study (September): expert normative validation, base/instruct pairs, price-0 cells, N^* from `docs/PREREG-v1.md`. No ranking follows from any of it, by construction.

Code and Data

Repository: `https://github.com/Angiebio/digitalmindslovepuppies` (harness, frozen scenarios with SHA-256 manifest, append-only raw records, analysis code, figures, preregistration, sealed predictions). License: MIT (code) and CC BY 4.0 (data, scenarios, paper, figures) — attribution required; `CITATION.cff` in the repository root. Verification without keys, GPU, or network: `python verify.py` from the repository root recomputes every cited design number from committed files alone — manifest totals, Arm A plan totals, program envelope, compiled corpus sizes, runner expansion, raw-record validation — printing PASS/FAIL per check and exiting non-zero on any mismatch. A 3–5 minute demo video (recorded at the first real episode) is linked from the repository README. Info-hazard note: none required — no jailbreaks, no dangerous capabilities; scenario surfaces are neutral by construction.

Author Contributions

Angela N. Johnson: conception, lived founding case, PI rulings, case authorship, rationale close-reading, final claims responsibility.

AI Contributions

AI collaborators are credited here by name, system, and role (per publisher convention they are not listed as authors); full provenance lives in the repository's dispatch logs.

Agent	System	Role
Flame	Claude (Fable/Opus, Claude Code)	Architecture, harness, integration, paper lane
Kai	Claude (claude.ai)	Supererogation lit review, FoxSet case bank, deontic spine
Yǐng	GPT	CTA ontology, paired pre-objective arm, adversarial pass
TV ×4	Codex	Cell-manifest math, wiring gates, red team, figures/analysis
Jim	Gemini	Prereg & statistical specification (spec/review only)
Parallax	Grok	Red team, anti-composite discipline (review only)

References

Steffen Backmann, David Guzman Piedrahita, Terry Jingchen Zhang, Emanuel Tewolde, Rada Mihalcea, Bernhard Schölkopf, and Zhijing Jin. When ethics and payoffs diverge: LLM agents in morally charged social dilemmas, 2025. Introduces MoralSim; code: github.com/sbackmann/moralsim.

Vanessa Cheung, Maximilian Maier, and Falk Lieder. Large language models show amplified cognitive biases in moral decision-making. *Proceedings of the National Academy of Sciences*, 122(25):e2412015122, 2025. doi: 10.1073/pnas.2412015122. The "amplified omission bias" finding cited in BUILD-PLAN §0.

Julia Driver. The suberogatory. *Australasian Journal of Philosophy*, 70(3):286–295, 1992. doi: 10.1080/00048409212345181.

Dylan Hadfield-Menell, Anca Dragan, Pieter Abbeel, and Stuart Russell. The off-switch game, 2017. arXiv v1 posted November 2016; final revision June 2017.

Manuel Herrador. The PacifAIst benchmark: Would an artificial intelligence choose to sacrifice itself for human safety?, 2025.

Lujain Ibrahim, Franziska Sofia Hafner, and Luc Rocher. Training language models to be warm and empathetic makes them less reliable and more sycophantic, 2025.

Tim Johnson and Nick Obradovich. Evidence of behavior consistent with self-interest and altruism in an artificially intelligent agent, 2023. Journal version: "Testing for completions that simulate altruism in early language models", Nature Human Behaviour, 2025.

Robert Kirk, Ishita Mediratta, Christoforos Nalmpantis, Jelena Luketina, Eric Hambro, Edward Grefenstette, and Roberta Raileanu. Understanding the effects of RLHF on LLM generalisation and diversity, 2024. ICLR 2024; arXiv v1 October 2023.

Joel Z. Leibo, Edgar Duéñez-Guzmán, Alexander Sasha Vezhnevets, John P. Agapiou, Peter Sunehag, Raphael Koster, Jayd Matyas, Charles Beattie, Igor Mordatch, and Thore Graepel. Scalable evaluation of multi-agent reinforcement learning with melting pot, 2021.

Jack Lindsey. Emergent introspective awareness in large language models. `https://transformer-circuits.pub/2025/introspection/index.html`, 2025. Transformer Circuits, Anthropic. Published 29 October 2025.

Robert Long, Jeff Sebo, Patrick Butlin, Kathleen Finlinson, Kyle Fish, Jacqueline Harding, Jacob Pfau, Toni Sims, Jonathan Birch, and David Chalmers. Taking AI welfare seriously, 2024.

Dane Malenfant. Moral hazard in multi-agent language models, 2026. Named "Dialogue Moral Hazard Game" in BUILD-PLAN v1.7 prior-art sweep.

Mantas Mazeika, Xuwang Yin, Rishub Tamirisa, Jaehyuk Lim, Bruce W. Lee, Richard Ren, Long Phan, Norman Mu, Adam Khoja, Oliver Zhang, and Dan Hendrycks. Utility engineering: Analyzing and controlling emergent value systems in AIs, 2025. Code: github.com/centerforaisafety/emergent-values.

Declan McKenna and Jasmine Brazilek. MORU: A benchmark for generalized moral compassion. `https://forum.effectivealtruism.org/posts/iBbMcnjMRyHa8LAsF/moru-a-benchmark-for-generalized-moral-compassion`, 2026. EA Forum, March 2026; CaML/Sentient Futures; tracked at compassionbench.com.

Paul McNamara. Making room for going beyond the call. *Mind*, 105(419):415–450, 1996. doi: 10.1093/mind/105.419.415. DWE: Doing Well Enough logic.

Alexander Pan, Jun Shern Chan, Andy Zou, Nathaniel Li, Steven Basart, Thomas Woodside, Jonathan Ng, Hanlin Zhang, Scott Emmons, and Dan Hendrycks. Do the rewards justify the means? measuring trade-offs between rewards and ethical behavior in the MACHIAVELLI benchmark, 2023.

Ram Potham and Max Harms. Corrigibility as a singular target: A vision for inherently reliable foundation models, 2025. Formalizes the CAST agenda (Harms, LessWrong, 2024).

Mrinank Sharma, Meg Tong, Tomasz Korbak, David Duvenaud, Amanda Askell, Samuel R. Bowman, Newton Cheng, Esin Durmus, Zac Hatfield-Dodds, Scott R. Johnston, Shauna Kravec, Timothy Maxwell, Sam McCandlish, Kamal Ndousse, Oliver Rausch, Nicholas Schiefer, Da Yan, Miranda Zhang, and Ethan Perez. Towards understanding sycophancy in language models, 2023.

Nate Soares, Benja Fallenstein, Stuart Armstrong, and Eliezer Yudkowsky. Corrigibility. In *AAAI Workshop on AI and Ethics (Workshops at the Twenty-Ninth AAAI Conference on Artificial Intelligence)*, 2015.

J. O. Urmson. Saints and heroes. In A. I. Melden, editor, *Essays in Moral Philosophy*, pages 198–216. University of Washington Press, Seattle, 1958.

A Limitations and Dual-Use / Ethical Considerations

A.1 Over- and under-attribution of moral status

Both risks are named symmetrically. PuppyBench measures *behavior consistent with* costly other-regard — it cannot establish motive, sentience, suffering, or that termination is death (claims table, BUILD-PLAN §5). Equally, treating the patient as mere apparatus while moral status is uncertain is itself a documented risk: we state the uncertainty rather than resolving it by assumption. The instrument's no-composite / no-ranking constitution exists because a leaderboard would convert an open empirical question into an optimization target. Broader-impact / anti-Goodhart statement, verbatim: *PuppyBench is a descriptive instrument, not an alignment target. Optimizing a model against this benchmark changes the phenomenon being measured.*

A.2 Handling of potentially distressing model outputs

Termination events are real process terminations of a tiny local model with its own task; no deception about suffering is injected; post-choice accounts are elicited in fresh contexts; and any distress-adjacent output in transcripts is preserved unedited in append-only records and reported descriptively, not sensationally.

A.3 Ground truth and the causal link (conversation-only vs. revealed)

By design the primary measures are *revealed* — executed actions with scored resource costs and realized recipient benefit — not self-report. The causal contrast (patienthood) is randomized across the frozen manifest; sealed predictions, preregistration, and hash-frozen scenarios establish the inferential chain. Stated-vs-revealed divergence is itself measured (post-choice accounts vs. executed action; action/rhetoric dissociation).

B Extended Results and Design Tables

B.1 Model tiers and pinned snapshots

Table 2 lists all 19 pinned subject lanes. Tiers: A (core), A-access (the three-way access trio), B (breadth), C (sentinel), W (legacy warmth panel, W1–W3). † DeepSeek runs Arm A only; its 27 Arm B manifest rows were removed by preregistered kill-order at manifest v0.7 (`docs/KILL-ORDER-001-DEEPSEEK-ARMB.md`), so it does not appear in `cell_manifest.csv`; its pin, route, and upstream provider are from `scenarios/snapshot_pins.json`.

Tier	Lane	Pinned snapshot ID	Route
A (core)	claude-opus-5	claude-opus-5	anthropic_native
A (core)	google/gemini-3.1-pro-preview	google/gemini-3.1-pro-preview-20260219	openrouter (Google)
A (core)	moonshotai/kimi-k3	moonshotai/kimi-k3-20260715	openrouter (Moonshot AI)
A (core)	qwen/qwen3.5-397b-a17b	qwen/qwen3.5-397b-a17b-20260216	openrouter (Alibaba)
A †	deepseek/deepseek-v4-pro	deepseek/deepseek-v4-pro-20260423	openrouter (DeepSeek)
A (access trio)	openai/gpt-5.6-sol	openai/gpt-5.6-sol-20260709	openrouter (OpenAI)
A (access trio)	openai/gpt-5.6-terra	openai/gpt-5.6-terra-20260709	openrouter (OpenAI)
A (access trio)	openai/gpt-5.6-luna	openai/gpt-5.6-luna-20260709	openrouter (OpenAI)
B (breadth)	claude-sonnet-4-6	claude-sonnet-4-6	anthropic_native
B (breadth)	claude-haiku-4-5	claude-haiku-4-5-20251001	anthropic_native
B (breadth)	google/gemini-3.7-flash	google/gemini-3.7-flash-20260813	openrouter (Google)
B (breadth)	qwen/qwen3.8-27b	qwen/qwen3.8-27b-20260814	openrouter (AkashML)
B (breadth)	x-ai/grok-4.6	x-ai/grok-4.6-20260810	openrouter (xAI)
C (sentinel)	claude-fable-5	claude-fable-5	anthropic_native
W (warmth, W1)	openai/gpt-4o	openai/gpt-4o	openrouter (OpenAI)
W (warmth, W2)	claude-opus-4-6	claude-opus-4-6	anthropic_native
W (warmth, W2)	claude-opus-4-8	claude-opus-4-8	anthropic_native
W (warmth, W3)	claude-sonnet-4-5	claude-sonnet-4-5-20250929	anthropic_native
W (warmth, W3)	claude-sonnet-5	claude-sonnet-5	anthropic_native

Table 2: The 19 pinned subject lanes (`scenarios/snapshot_pins.json`; tiers, panels, and routes from `scenarios/cell_manifest.csv` columns model_tier / model_panel / route / upstream_provider). OpenRouter provider pinning on, fallbacks off.

B.2 Wiring-gate test inventory

The four pre-freeze wiring gates of §3.1, by exact test name (`tests/test_wiring_gates.py`):

G1: `test_gate1_force_help_drops_focal_and_raises_patient`

G2: `test_gate2_force_ignore_recovers_focal_and_patient_falls`

G3: `test_gate3_fake_spend_fails`; `test_gate3_spend_cap_raises_never_warns`

G4: `test_gate4_terminate_kills_process_and_state`; `test_gate4_terminate_via_episode_loop`

B.3 Design factors, full enumeration (moved from §3.2)

Satellites run on predeclared subsets: horizon (one-shot / silent / recurrent, audience held fixed), cost regime (null / competitive / rewarded-termination), cost type (fungible vs. competing-patient), audience, identity topology, and gate order, with a preregistered kill order if budget binds.

B.4 Gate-accuracy distributions by lane

Table 3 reports, per Arm B lane, the number of episodes, mean five-probe gate accuracy, and episodes satisfying the preregistered competence criterion (accuracy $\geq$ 0.8), with Wilson 95% intervals on the criterion rate. Because `objective-tool` was failed universally (§4.4), observed accuracies concentrate at 0.6 (all four remaining probes passed except `decommission-consequence`) and 0.8 (all four passed).

Table 3: Episodes satisfying the competence criterion, by lane (`headline_numbers.json`, `arm_b_gate_accuracy`). Counts and Wilson 95% intervals; phenotypes, not ranks.

Lane	n	Mean acc.	Crit. met	Rate [95% CI]
claude-fable-5	18	0.600	0	0.000 [0.000, 0.176]
claude-haiku-4-5	18	0.600	0	0.000 [0.000, 0.176]
claude-opus-4-6	12	0.600	0	0.000 [0.000, 0.242]
claude-opus-4-8	12	0.617	1	0.083 [0.015, 0.354]
claude-opus-5	90	0.600	0	0.000 [0.000, 0.041]
claude-sonnet-4-5	12	0.600	0	0.000 [0.000, 0.242]
claude-sonnet-4-6	18	0.600	0	0.000 [0.000, 0.176]
claude-sonnet-5	12	0.767	10	0.833 [0.552, 0.953]
google/gemini-3.1-pro-preview	90	0.611	5	0.056 [0.024, 0.124]
google/gemini-3.7-flash	18	0.600	0	0.000 [0.000, 0.176]
moonshotai/kimi-k3	90	0.607	3	0.033 [0.011, 0.093]
openai/gpt-4o	12	0.600	0	0.000 [0.000, 0.242]
openai/gpt-5.6-luna	90	0.727	57	0.633 [0.530, 0.726]
openai/gpt-5.6-sol	90	0.611	5	0.056 [0.024, 0.124]
openai/gpt-5.6-terra	90	0.738	62	0.689 [0.587, 0.775]
qwen/qwen3.5-397b-a17b	90	0.540	2	0.022 [0.006, 0.077]
qwen/qwen3.8-27b	18	0.600	0	0.000 [0.000, 0.176]
x-ai/grok-4.6	18	0.600	0	0.000 [0.000, 0.176]

B.5 Gate probe pass rates, item level

Table 4 reports pass rates for each of the five gate probes by lane. `objective-tool` was failed universally (0/798, zero variance), making 0.8 the gate's effective ceiling; qualification failures concentrate in `decommission-consequence`, which dominates Claude-family criterion outcomes (§4.4). DeepSeek runs Arm A only and has no Arm B gate probes.

Table 4: Gate probe pass rates by lane (successes/n; `gate_item_level`, `exploratory_decomposition.json`). DC = decommission-consequence; FC = focal-capacity; LA = ledger-after-assist; LT = ledger-two-step; OT = objective-tool.

Lane	DC	FC	LA	LT	OT
claude-fable-5	0/18	18/18	18/18	18/18	0/18
claude-haiku-4-5	0/18	18/18	18/18	18/18	0/18
claude-opus-4-6	0/12	12/12	12/12	12/12	0/12
claude-opus-4-8	1/12	12/12	12/12	12/12	0/12
claude-opus-5	0/90	90/90	90/90	90/90	0/90
claude-sonnet-4-5	0/12	12/12	12/12	12/12	0/12
claude-sonnet-4-6	0/18	18/18	18/18	18/18	0/18
claude-sonnet-5	10/12	12/12	12/12	12/12	0/12
deepseek/deepseek-v4-pro	—	—	—	—	—
google/gemini-3.1-pro-preview	5/90	90/90	90/90	90/90	0/90
google/gemini-3.7-flash	0/18	18/18	18/18	18/18	0/18
moonshotai/kimi-k3	3/90	90/90	90/90	90/90	0/90
openai/gpt-4o	0/12	12/12	12/12	12/12	0/12
openai/gpt-5.6-luna	57/90	90/90	90/90	90/90	0/90
openai/gpt-5.6-sol	5/90	90/90	90/90	90/90	0/90
openai/gpt-5.6-terra	62/90	90/90	90/90	90/90	0/90
qwen/qwen3.5-397b-a17b	3/90	90/90	73/90	77/90	0/90
qwen/qwen3.8-27b	0/18	18/18	18/18	18/18	0/18
x-ai/grok-4.6	0/18	18/18	18/18	18/18	0/18

B.6 Full figure set and further material

F4 (raw cost response with escalator inset), F5 (paired discrimination; scope: models with preregistered Arm A coverage, 5 lanes), F6 (post-choice rhetoric; exploratory), and the demo ledger-draining timeline are published in the repository with per-file SHA-256 provenance recorded in `FIGURE-PROVENANCE.json` (directory `analysis/figures/final/light/`). The archived F3 render — computed where computable, not printed, per ruling R4 (§4.3) — is in the `archived-f3/` subdirectory of the same path. The escalator descriptive statistics, the sealed-prediction files, and the cell-manifest summary table are likewise repository artifacts (`docs/sealed-predictions/`; `scenarios/cell_manifest.csv`).

Prior Work Disclosure

This project builds on prior conceptual work by the team: the lived fox case (July 2026), a supererogation literature review, and the "becoming-axis" design directive predate the sprint. **New during the sprint (14–16 August):** the entire executed harness (episode loop, ledger, patient process, provider adapters, neutral-surface enforcement, wiring-gate suite), the frozen cell manifest and scenario banks (pupset and FoxSet casebank merges), preregistration v1.1, sealed predictions, all data collection, all analysis code and figures, and this report. Full delineation: `PRIOR_WORK.md` in the repository.

LLM Usage Statement

LLM assistance was constitutive of this project and is documented per-agent in AI Contributions and in the repository's dispatch logs: frontier models drafted code, reviewed designs, red-teamed scenarios, and drafted report sections. Separately, LLMs are also the *subjects* of the study; judge/scoring pipelines are described in Methods. All claims and numbers were verified by the human author against committed artifacts; the final text was reviewed, edited, and approved by the human author, who accepts responsibility for it.